

Week 5 Project: IS607

Alexander Satz

October 3, 2014

A data set was downloaded from <https://www.ebi.ac.uk/chembl/sarfari/kinasesarfari/>. This data set contains 30,016 potency measurements of compounds against various protein kinases. The data set derives from thousand of publications and ChEMBL internal sources.

I sought to answer the following question: Do compounds with high potency in some targets end up being assayed more often, i.e. are 'potent' cmpds followed-up?

1. Open the file

```
> library("dplyr")
> library("tidyr")
> library("ggvis")
> file <- read.table(file = "/Users/alexandersatz/Documents/Cuny/IS607/week5/ks_bioactivity.")
> bioact.1 <- tbl_df(file)
```

2. As shown below, there are 1447 types of measurements. This needs to be cleaned up before we can classify potency. Some values are log, some -log and some are not logs at all. Each row is an observation, and so there is no need for 'pivoting' etcetera.

```
> summarise(bioact.1,
+           numberdatatype = n_distinct(ACTIVITY_TYPE))
```

Source: local data frame [1 x 1]

	numberdatatype
1	1447

```
> activity.type <- group_by(bioact.1, ACTIVITY_TYPE)
> types.df <- summarise(activity.type,
+                       number = n())
> types.df
```

Source: local data frame [1,447 x 2]

ACTIVITY_TYPE	number
---------------	--------

```

1  -Delta G obs      2
2    -Log alpha      1
3      -Log C         7
4    -Log EC50       21
5    -Log IC25        1
6    -Log IC50        7
7  -Log IC50(M)       4
8      -Log KB        1
9      -Log KD        3
10     -Log KD50       4
..      ...      ...

```

Additionally there are numerous scales being used including nM and uM, stated in both lower and uppercase (see below). Before all values can be converted to log units, this will need to be standardized.

```

> activity.units <- group_by(bioact.1, STANDARD_UNIT)
> units.df <- summarise(activity.units,
+                        number = n())
> units.df <- data.frame(units.df)
> head(units.df)

```

	STANDARD_UNIT	number
1		47944
2	(mg of CPT) kg-1	1
3	(nM of XMP formed) hr-1 (mg of protein)-1	2
4	(ug of base) ml-1	4
5	(ug of cross-linked protein) (mg of protein)-1	3
6	/hr	3

First we tackle those values already present on a log10 scale. We run a grepl search for 'log'. The result includes nonsensical outliers and measurements such as 'logD' and 'logP' which are not activity measurements. These rows are removed by additional filter() using text matching. Lastly, the Log value is converted to an integer because we want there to be a limited number of potency 'levels'. The data frame still has many columns as we haven't decided what to get rid of yet. NA values also exist. NA values may derive from assays where a 'value' for that particular compound could not be calculated. These values should not be in the database. I will remove them later.

```

> bioact.log <- filter(bioact.1, grepl("log", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.log <- filter(bioact.log, ! grepl("log2", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.log <- filter(bioact.log, ! grepl("logp", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.log <- filter(bioact.log, ! grepl("logd", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.log <- filter(bioact.log, ! grepl("GI50", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.logged1 <- mutate(bioact.log,
+                          LOG.ACT = as.integer(abs(STANDARD_VALUE)))

```

```
> bioact.logged1 <-arrange(bioact.logged1, desc(LOG.ACT))
> bioact.na <- (bioact.logged1[is.na(bioact.logged1$LOG.ACT),])
> bioact.logged1
```

Source: local data frame [2,317 x 14]

	ACTIVITY_ID	DOM_ID	NAME	ASSAY_TYPE	COMPOUND_ID	ACTIVITY_TYPE
1	2582768	1553	hEGFR_1553	B	35820	Log IC50
2	2711279	NA	Starlite ADMET	A	374400	log KOA
3	2582767	1553	hEGFR_1553	B	328106	Log IC50
4	2383811	1553	hEGFR_1553	B	271122	Log IC50
5	2383814	1553	hEGFR_1553	B	299622	Log IC50
6	2383291	1553	hEGFR_1553	B	294475	Log IC50
7	2383328	1553	hEGFR_1553	B	328216	Log IC50
8	2437590	NA	Starlite ADMET	A	229760	log KOA
9	2436949	NA	Starlite Functional	F	464859	Log EC50
10	2446933	NA	Starlite ADMET	A	415	Log k'
..

Variables not shown: RELATION (fctr), STANDARD_VALUE (dbl), STANDARD_UNIT (fctr), ACTIVITY_COMMENT (fctr), ChEMBL_ACTIVITY_ID (int), ChEMBL_ASSAY_ID (int), PUBMED_ID (int), LOG.ACT (int)

>

Next I need to deal with 'values' measured not on a log scale. These can have either nM or uM scales. First I pull out all values that are NOT 'log' values via a grepl match, then I divide this data into those that are on the nM and uM scales.

```
> bioact.notlog <-filter(bioact.1, ! agrepl("log", ACTIVITY_TYPE, ignore.case = TRUE))
> bioact.loguM <-bioact.notlog %>%
+   filter(grepl("um", STANDARD_UNIT, ignore.case = TRUE)) %>%
+   mutate(LOG.ACT = as.integer(-log10((STANDARD_VALUE/1000000))))
> bioact.loguM <-arrange(bioact.loguM, desc(LOG.ACT))
> #bioact.loguM$LOG.ACT ## looks great and values range from 4-6 mainly, so the right range
>
> bioact.lognM <-bioact.notlog %>%
+   filter(grepl("nm", STANDARD_UNIT, ignore.case = TRUE)) %>%
+   mutate(LOG.ACT = as.integer(-log10((STANDARD_VALUE/1000000000))))
> bioact.lognM <-arrange(bioact.lognM, desc(LOG.ACT))
> #bioact.lognM$LOG.ACT ## looks good
```

Now combine the 3 dataframes. I have 11000 entries. The final product can be inspected to see that the calculated value LOG.ACT matches the expected value! Last, NA values are removed as we know from above inspection that they are 'garbage' in the data set.

```

> biact2 <- rbind(bioact.logged1, bioact.loguM, bioact.lognM)
> biact2 <- select(biact2, COMPOUND_ID, STANDARD_UNIT, STANDARD_VALUE, LOG.ACT, DOM_ID, NAME)
> biact2 <- arrange(biact2, desc(LOG.ACT))
> head(biact2, 10)

```

Source: local data frame [10 x 6]

	COMPOUND_ID	STANDARD_UNIT	STANDARD_VALUE	LOG.ACT	DOM_ID	NAME
1	59	nM	8.00e-06	14	NA	Starlite Functional
2	1532440	uM	4.07e-07	12	NA	Starlite Functional
3	1377737	uM	1.36e-07	12	NA	Starlite Functional
4	1494120	uM	4.07e-07	12	NA	Starlite Functional
5	1458022	uM	4.07e-07	12	NA	Starlite Functional
6	216933	nM	8.80e-04	12	1950	hABL1_1950
7	408247	nM	7.30e-04	12	1950	hABL1_1950
8	223228	nM	2.50e-04	12	NA	Starlite Functional
9	223228	nM	3.00e-04	12	NA	Starlite Functional
10	264189	nM	6.60e-04	12	1950	hABL1_1950

```

> biact3 <- biact2[!is.na(biact2$LOG.ACT),]

```

3. Now we group by 'compound id' and determine the max potency of each compd in any assay.

```


> biact.cmp <- group_by(biact3, COMPOUND_ID)
> bioact.sum <- summarize(biact.cmp,
+                         max.pot = max(LOG.ACT),
+                         tot.Assays = n(),
+                         uniq.Assays = n_distinct(NAME)
+                         )
> bioact.sum <- arrange(bioact.sum, desc(max.pot))
> bioact.sum

```

Source: local data frame [49,475 x 4]

	COMPOUND_ID	max.pot	tot.Assays	uniq.Assays
1	59	14	83	4
2	76	12	1087	3
3	129	12	1311	4
4	135	12	125	2
5	633	12	73	3
6	1563	12	142	3
7	216933	12	2	1
8	223228	12	687	3
9	264189	12	5	1
10	406721	12	2	1
..

4. We plot potency of each cmpd (its max potency in any assay) versus number of the unique assays the compound was run in.




From the figure above it can be observed that some relatively potent cmpds (Log values of 7-10) have been tested in a large number of different assays (>300). These cmpds are likely 'standards' and there are relatively few of them. The smoothed line and transparent points available in the ggvis package more clearly show this, however the ggvis pkg appear to be incompatible with sweave.

5. We want to know if the median number of unique assays done increases with potency of the compounds. We use the median and not the average as the 'standards' will heavily influence the mean. The purpose of this is to determine if 'potent' cmpds are followed up? *Or is the kinase database more of a data dump?*

```
> biact.group <- group_by(bioact.sum, max.pot)
> uniq.assays<- summarise(biact.group,
+                           v.avg = mean(uniq.Assays),
+                           v.med = median(uniq.Assays))
```

A plot of binned potency versus median number of unique assays run is then



provided:

We see that cmpds with higher potencies do not seem to be consistently run in a variety of different assays. Indeed, most compounds are only tested in 2 different assays and this appears independent of potency level. Generally, the important range of log potency is between 5 and 9. This area of the plot is flat.